

Daniel Aaron Levenstein

| [LinkedIn Profile](#) | [GitHub Portfolio](#) | [Kaggle Notebooks](#) | [Deployed AI Models](#) | 512-885-5925 |

Data Engineer

Data Engineer with experience building and validating distributed data pipelines in financial and healthcare domains. Proficient in PySpark, SQL, and AWS, with a focus on data integrity, large-scale data processing, and pipeline reliability. Background includes implementing data quality frameworks, developing reprocessing workflows, and supporting production-grade data systems handling high-volume datasets.

Certifications

- AWS Cloud Practitioner, AWS (2025)
- AI and ML Certification, Great Learning (2026)

Technical Skills

- Data Engineering & Analytics: Apache Spark, SQL, Pandas, NumPy, Apache Airflow
- Machine Learning & AI: Scikit-Learn, OpenCV, embeddings, LLMs, RAG systems
- Cloud & Infrastructure: AWS Lambda, Step Functions, Batch, ECS/EKS, S3, RDS, CloudWatch, CodeBuild
- Databases & Storage: SQL, DynamoDB, Vector databases

Professional Experience

Freelance Data Engineer

Austin, TX

| Self Employed

| 01/2025 - Present

Built data processing and ML pipelines using Python and Scala, including dataset ingestion, transformation, and reporting. (Python, Scala)

- [AWS Documentation RAG](#): Built a RAG app for querying AWS docs: crawled and preprocessed pages, indexed embeddings in a vector DB, and served LLM-powered technical answers via Streamlit.
- [HelmNet - Data Quality Case Study](#): Built image classification CNN models from different source datasets and generated a cross-validation report across all models using a holdout dataset. Increased final accuracy metric by 17%.
- [Prompt Engineering Case Study](#): Build a medical assistant chatbot using a publicly available instance of Llama model and evaluated the model on answer accuracy and instruction fidelity.
- [Stress vs. Physical Activity Correlation](#): Analyzed relationships between stress levels and physical activity using Kaggle dataset.

Software Developer

Rockville, MD

| Finra

| 10/2022 – 12/2024

Developed and validated scalable ETL pipelines using Apache Spark to ensure accurate data ingestion. (Java, Scala, AWS)

- Designed validation logic to detect anomalies and enforce consistency across large-scale financial data feeds.
- Generated test data for AWS Step Functions using Apache Spark and Hibernate, improving ingestion workflows.
- Developed and validated scalable Spark-based ETL pipelines supporting analytical workloads over large-scale datasets.

Pipeline Engineer

Austin, TX

| b.well Connected Health

| 09/2020 – 09/2022

Monitored the integrity of a healthcare app using a Test-driven development methodology. (PyTest/ Selenium).

- Designed and maintained AWS-based infrastructure supporting automated data and API validation workflows.
- Developed serverless monitoring solutions using AWS Lambda to track pipeline health and detect failures in near real-time.
- Automated validation processes for healthcare application data, reducing manual verification effort and improving reliability.
- Contributed to CI/CD pipelines to ensure consistent deployment and validation of data processing systems.

QA Engineer

Austin, TX

| Apiture

| 02/2019 – 08/2020

Conducted integration testing to validate interactions between backend APIs and front-end sites. (Java/ Node)

- Utilized PostgreSQL to create, update, and manage mock datasets, ensuring comprehensive test coverage.

Sr. QA Engineer

Austin, TX

| Conduent

| 06/2011 – 01/2019

Ran the QA team single-handed while the company was going through an acquisition. (Java)

Education

- Post Graduate Certificate in AI and Machine Learning, University of Texas (2026)
- Bachelor of Computer Science, Texas State University (2011)